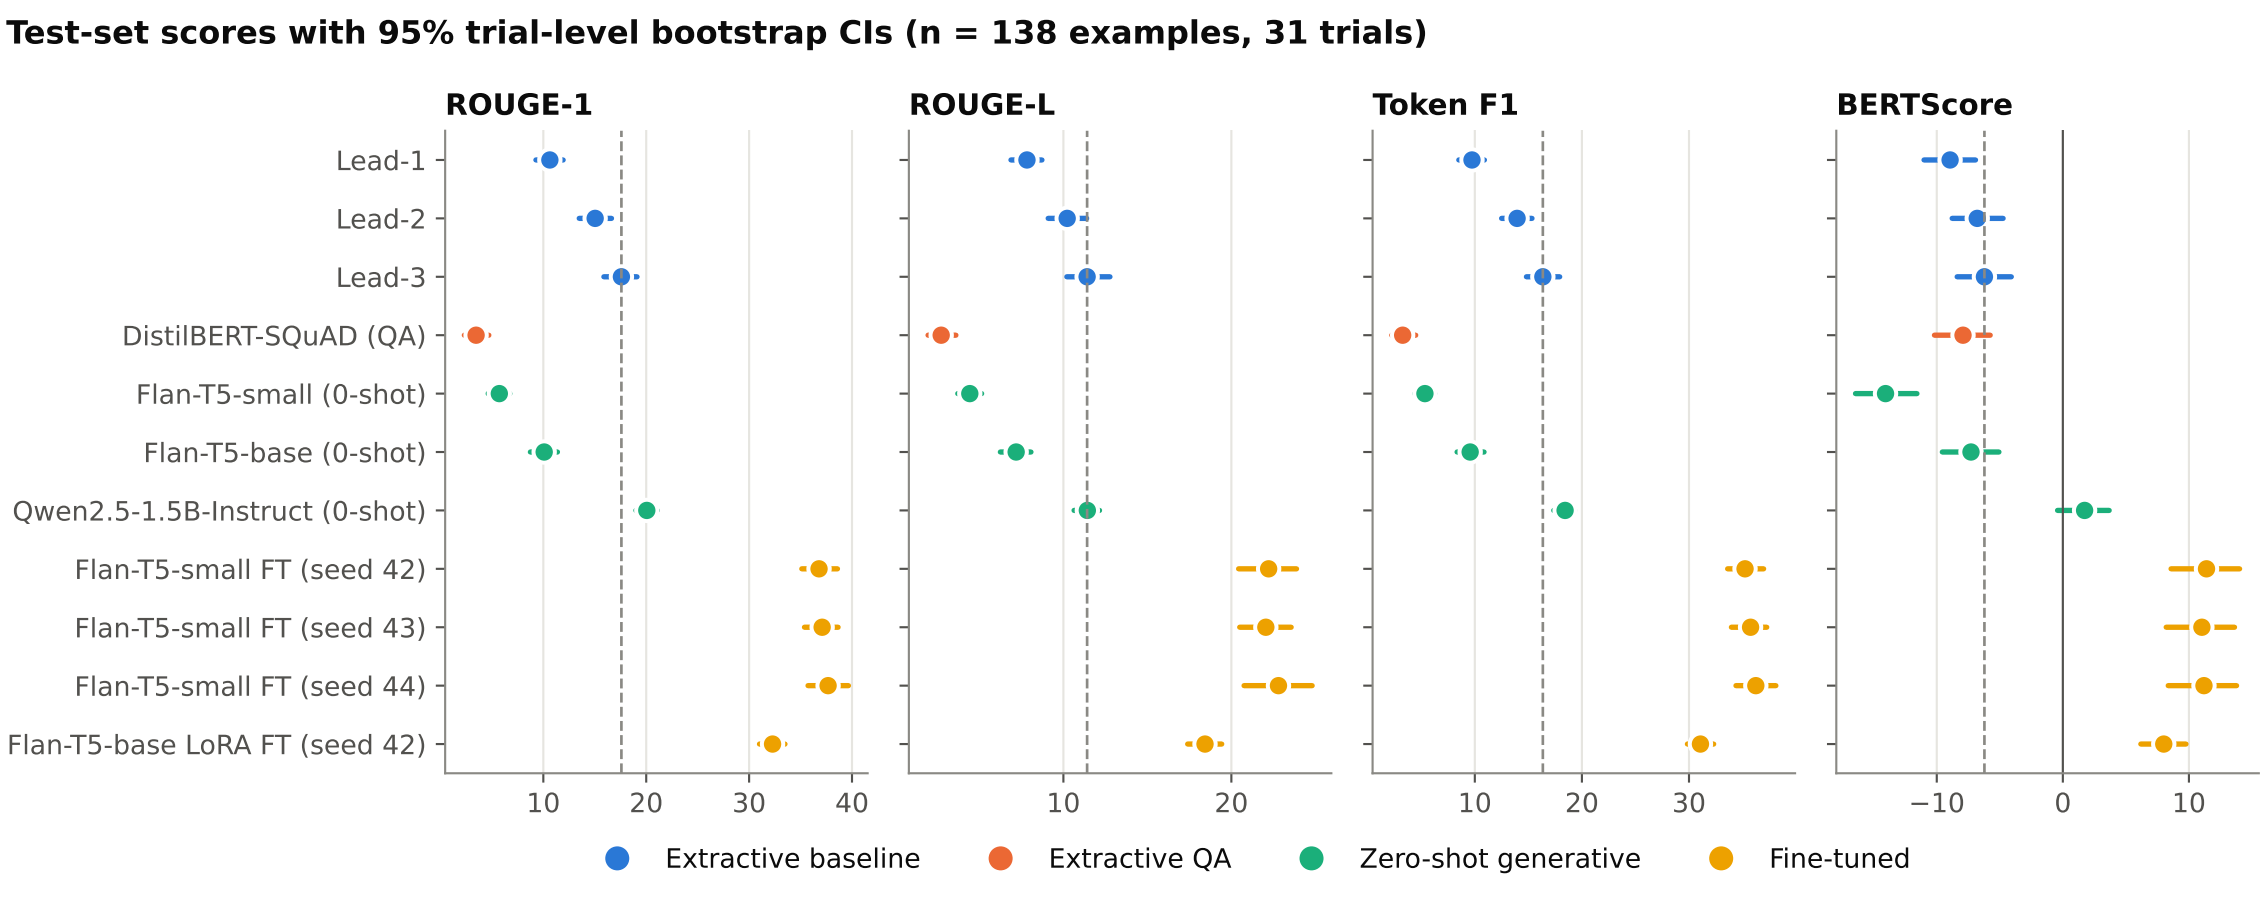
Test-set scores with 95% trial-level bootstrap CIs (n = 138 examples, 31 trials)
ROUGE-1
ROUGE-L
Token F1
BERTScore
Lead-1
Lead-2
Lead-3
DistilBERT-SQuAD (QA)
Flan-T5-small (0-shot)
Flan-T5-base (0-shot)
Qwen2.5-1.5B-Instruct (0-shot)
Flan-T5-small FT (seed 42)
Flan-T5-small FT (seed 43)
Flan-T5-small FT (seed 44)
Flan-T5-base LoRA FT (seed 42)
10
20
30
40
10
20
10
20
30
−10
0
10
Extractive baseline
Extractive QA
Zero-shot generative
Fine-tuned